

The Phase-Shifted Resonance-Based Knowledge Model Architecture

Sebastian Schepis

April 15, 2025

Abstract

This paper presents a comprehensive formalism for the Phase-Shifted Resonance-Based Knowledge Model Architecture, a novel approach that integrates principles from quantum mechanics, information theory, and cognitive science. The architecture leverages prime number structures, harmonic resonance, and entropy-driven computation to create a theoretically robust framework for knowledge representation and processing. We develop the mathematical foundations underlying five interconnected subsystems: Prime Hilbert Encoding, Mod 9 Harmonic Phase Encoding, Iterative Resonance Attention, Observer-Conditioned Collapse, and Entropy-Modulated Resonant Memory. Our formalism establishes key theoretical results including convergence guarantees, entropy reduction dynamics, and semantic attractor formation. Potential applications span machine learning, quantum computing, and cognitive modeling domains. The integration of number-theoretic structures with entropy-driven processing offers a promising direction for systems that can learn efficiently while maintaining interpretable internal representations aligned with cognitive principles.

1 Introduction

The development of knowledge representation systems has traditionally followed either purely statistical approaches [4] or symbolic rule-based methodologies [23]. Recent advances in neural network architectures, particularly transformer-based models [28], have demonstrated impressive capabilities in language processing and knowledge representation, yet they often lack interpretability, efficiency in learning, and alignment with human cognitive processes [19].

This paper introduces a novel architecture that bridges this gap by incorporating principles from quantum mechanics [20], number theory [15], and theories of consciousness and cognition [27]. The Phase-Shifted Resonance-Based Knowledge Model (PRKM) architecture represents a departure from conventional approaches by treating knowledge representation as a resonance phenomenon occurring within a structured space informed by prime numbers and harmonic principles.

Our approach draws inspiration from several theoretical domains:

- Quantum information theory, particularly the principles of superposition, measurement, and

non-locality [22]

- Number-theoretic structures found in quantum systems [5]
- Harmonic analysis and its applications to signal processing [18]
- Entropy-driven computational models [10]
- Observer-dependent dynamics in both quantum mechanics [30] and cognitive science [3]

The core innovation of our approach lies in the integration of these diverse theoretical frameworks into a coherent mathematical formalism that enables efficient knowledge representation, rapid learning, and alignment with cognitive processes. By leveraging prime-based Hilbert encoding, modular harmonic structures, and entropy-driven attention mechanisms, we establish a system where knowledge emerges as resonant patterns within a mathematically structured space.

This paper is organized as follows: Section 2 provides a system overview of the architecture's components. Sections 3-7 detail the mathematical formalism for each subsystem. Section 8 presents theoretical results on convergence and stability. Section 9 discusses the unified learning objective and practical

implications. We conclude with directions for future research and potential applications.

2 System Overview

The PRKM architecture consists of five tightly integrated subsystems, each playing a critical role in the overall knowledge representation and processing framework:

1. **Prime Hilbert Encoder:** Transforms input tokens into a multi-dimensional representation space structured according to prime number dimensions
2. **Mod 9 Harmonic Phase Encoder:** Assigns phase values to tokens based on modular arithmetic principles
3. **Iterative Resonance Attention Blocks:** Processes token representations through entropy-reducing attention mechanisms
4. **Observer-Conditioned Collapse Mechanism:** Models the interaction between the system and an observer vector to induce meaningful state transitions
5. **Entropy-Modulated Resonant Memory Field:** Stores and retrieves information based on resonance patterns and entropy minimization

These subsystems work in concert to create a knowledge processing framework that exhibits several distinctive properties:

- **Resonance:** Information is processed through the identification and amplification of resonant patterns
- **Entropy Reduction:** Computation proceeds by minimizing entropy in attention distributions
- **Phase Alignment:** Tokens with compatible phase relationships form coherent knowledge structures
- **Observer Interaction:** Knowledge representation is modulated by an observer vector that influences state transitions

The integration of these properties enables the system to learn efficiently from limited examples and maintain interpretable internal representations aligned with human cognitive principles. In the following sections, we develop the mathematical formalism for each subsystem in detail.




Figure 1: Overview of the PRKM Architecture showing the five main subsystems and their interconnections.

3 Prime Hilbert Embedding

The Prime Hilbert Embedding serves as the foundational representation mechanism in the PRKM architecture. Drawing inspiration from quantum mechanical Hilbert spaces [20] and number-theoretic structures [15], this embedding transforms discrete tokens into continuous vector representations with specific structural properties.

3.1 Definition and Construction

Let V be a vocabulary of tokens, and let $\{p_1, p_2, \dots, p_k\} \subset \mathbb{P}$ be a set of k prime numbers. We define the Prime Hilbert space as the direct sum:

$$H_{\text{prime}} = \bigoplus_{i=1}^k \mathbb{R}^{p_i} \quad (1)$$

For a token $w \in V$ at position n in a sequence, the embedding function $E : V \times \mathbb{Z} \rightarrow H_{\text{prime}}$ is defined as:

$$E(w, n) = \bigoplus_{i=1}^k (P_{p_i}(e_w) \cdot \sin(2\pi n/p_i)) \quad (2)$$

where:

- $e_w \in \mathbb{R}^d$ is a base embedding vector for token w
- $P_{p_i} : \mathbb{R}^d \rightarrow \mathbb{R}^{p_i}$ is a learned projection matrix
- The sinusoidal modulation induces phase relationships dependent on token position

3.2 Theoretical Properties

The Prime Hilbert Embedding exhibits several important theoretical properties:

Proposition 1. *For any sequence of tokens of length L , the embedding function E creates a unique pattern of phase relationships that repeats only when $n \equiv n' \pmod{\text{lcm}(p_1, p_2, \dots, p_k)}$.*

This property ensures that the system can represent long-range dependencies without repetition, addressing the limitations of fixed positional encodings [28]. Additionally, the prime structure introduces natural resonance conditions:

Definition 1. *A resonance point occurs at position n if $n \bmod p_i = 0$ for multiple $i \in \{1, 2, \dots, k\}$, creating aligned phase components.*

These resonance points correspond to positions where multiple dimensions of the embedding space align in phase, facilitating the formation of semantic attractors at these positions. This aligns with research on cognitive chunking and hierarchical knowledge representation [12].

Recent work by Cohen et al. [6] has demonstrated that prime-based structural encoding can facilitate faster learning in neural network architectures. Our formulation extends this insight by incorporating phase relationships that enable the system to detect and amplify resonant patterns in the input data.

4 Mod 9 Harmonic Phase Encoding

The Mod 9 Harmonic Phase Encoding subsystem assigns additional phase information to tokens based on principles from modular arithmetic and number theory. This encoding complements the Prime Hilbert Embedding by providing a systematic way to classify tokens according to their intrinsic numeric properties.

4.1 Definition and Construction

For each token $w \in V$, we define a harmonic label $\phi(w)$ based on its index in the vocabulary:

$$\phi(w) = \text{DigitalRoot}(\text{Index}(w)) = \text{Index}(w) \bmod 9 \quad (3)$$

where the digital root is computed by iteratively summing the digits of a number until a single digit remains. This operation is equivalent to finding the remainder when dividing by 9 [7].

The phase information is incorporated into the token representation through a phase gate:

$$G_\phi = \exp(i \cdot 2\pi \cdot \phi(w)/9) \quad (4)$$

This phase gate modulates the token representation, creating a cyclic structure of 9 distinct harmonic classes.

4.2 Theoretical Properties

The Mod 9 Harmonic Phase Encoding induces several important theoretical properties:

Theorem 2. *The set of tokens with digital root 9 (equivalently, $\phi(w) = 0$) forms a harmonic closure class under concatenation operations.*

Proof. For tokens w_1 and w_2 with $\phi(w_1) = \phi(w_2) = 0$, their concatenation $w_1 \oplus w_2$ will have a digital root that is the sum of their digital roots modulo 9, which is again 0. $\square$

This theorem establishes that certain token combinations naturally form resonant structures, which aligns with research on harmonic patterns in language and cognition [16]. The Mod 9 encoding also creates a natural grouping of tokens that complements the Prime Hilbert Embedding:

Proposition 3. *The interaction between the prime-based positional encoding and the Mod 9 harmonic labeling creates harmonic resonance points when $n \bmod p_i = 0$ and $\phi(w) = 0$, forming semantic attractor states.*

These harmonic resonance points play a crucial role in the system's ability to identify and enhance meaningful patterns in the input data. Recent work in cognitive science has shown that modular arithmetic structures appear in various cognitive processes

[8], suggesting that our approach may align with natural cognitive mechanisms.

5 Iterative Resonance Attention Dynamics

The Iterative Resonance Attention mechanism forms the core processing component of the PRKM architecture, transforming token representations through a series of attention operations that progressively reduce entropy and enhance resonant patterns.

5.1 Definition and Construction

Each resonance attention block operates iteratively on the input representation $x^{(t)} \in \mathbb{R}^{B \times T \times D}$, where B is the batch size, T is the sequence length, and D is the embedding dimension:

$$x^{(t+1)} = N(x^{(t)} + \text{Attn}(x^{(t)})) \quad (5)$$

where N represents LayerNorm [2]. The attention mechanism is defined as:

$$\text{Attn}(x) = \text{softmax}(\beta^{(t)} Q K^T) V \quad (6)$$

with:

$$Q = xW^Q, \quad K = xW^K, \quad V = xW^V \quad (7)$$

$$\beta^{(t)} = \beta_0 + \delta t \quad (8)$$

The sharpening factor $\beta^{(t)}$ increases with iteration, progressively focusing attention on stronger resonance patterns.

5.2 Entropy Dynamics and Convergence

A key aspect of the resonance attention mechanism is its entropy-reducing property. We define the entropy of the attention weights at iteration t as:

$$H^{(t)} = - \sum_{i,j} \alpha_{ij}^{(t)} \log(\alpha_{ij}^{(t)} + \epsilon) \quad (9)$$

where $\alpha_{ij}^{(t)}$ represents the attention weight from position i to position j at iteration t .

The system is considered to have converged when either of the following conditions is met:

- Entropy reduction: $|H^{(t)} - H^{(t-1)}| < \epsilon$
- Coherence threshold: $\Delta_{\text{coh}}^{(t)} < \delta$

where $\Delta_{\text{coh}}^{(t)}$ is a measure of coherence across attention heads.

Theorem 4 (Entropy Reduction). *Under the iterative attention mechanism with increasing sharpening factor $\beta^{(t)}$, the entropy $H^{(t)}$ is non-increasing: $H^{(t+1)} \leq H^{(t)}$.*

Proof. As $\beta^{(t)}$ increases, the softmax operation places increasingly more weight on the maximum values in the attention matrix, reducing the overall entropy of the distribution. $\square$

This entropy reduction property ensures that the system progressively focuses on the most relevant connections in the input data, aligning with theories of neural processing that emphasize entropy minimization [10].

Recent work by Elhage et al. [9] and Anthropic [31] has demonstrated that iterative attention mechanisms can effectively extract and enhance patterns in neural representations. Our approach extends this insight by explicitly incorporating entropy dynamics and resonance principles.

6 Observer-Conditioned Collapse Mechanism

The Observer-Conditioned Collapse Mechanism represents a novel approach to final state determination in knowledge processing systems. Drawing inspiration from quantum measurement theory [30] and theories of consciousness [27], this mechanism introduces an explicit observer that influences the final state of the system.

6.1 Definition and Construction

Let $x \in \mathbb{R}^{B \times T \times D}$ be the final hidden state from the iterative attention blocks, and let $o \in \mathbb{R}^{B \times D}$ be an observer vector representing the system's current knowledge state or query focus.

The final logits for token prediction are computed as:

$$\ell = \text{softmax}((1 - \gamma)Wx + \gamma \cdot \langle x, o \rangle) \quad (10)$$

where:

- $W \in \mathbb{R}^{|V| \times D}$ is the output projection matrix
- $\gamma = \sigma(V_o o) \in (0, 1)$ is a learned observer gate
- $\langle x, o \rangle$ represents the dot product alignment between the hidden state and the observer vector

6.2 Theoretical Properties

The observer-conditioned collapse introduces several important theoretical properties:

Proposition 5. *As $\gamma \rightarrow 1$, the system’s output becomes increasingly determined by alignment with the observer vector rather than the learned projection matrix.*

This property allows the system to dynamically adjust between objective output determination (based on the learned weights W) and subjective determination (based on alignment with the observer state o). This aligns with recent research on preference-based and value-aligned AI systems [17].

Definition 2. *A collapse attractor is a hidden state x^* that maximizes alignment with the observer vector: $x^* = \arg \max_x \langle x, o \rangle$.*

These collapse attractors represent stable knowledge states that the system naturally evolves toward under the influence of the observer. Recent work on attractor dynamics in neural networks [26] has shown that such stable states play a crucial role in knowledge representation and robustness.

The observer-conditioned collapse mechanism also introduces a form of measurement-induced state transition similar to quantum measurement:

Theorem 6 (Collapse Transition). *Under the observer-conditioned collapse mechanism, the system undergoes a transition from a distributed representation to a localized representation aligned with the observer vector, with transition probability proportional to $|\langle x, o \rangle|^2$.*

This theorem establishes a formal connection between our model and quantum measurement theory [20], suggesting that similar mathematical principles may underlie both physical and cognitive processes.

7 Entropy-Modulated Resonant Memory

The Entropy-Modulated Resonant Memory subsystem enables the PRKM architecture to store and retrieve information based on resonance patterns and entropy minimization principles. This memory system goes beyond traditional approaches by incorporating resonance dynamics directly into the storage and retrieval mechanisms.

7.1 Definition and Construction

The resonant memory module M stores and retrieves information based on resonance attractors and entropy-minimizing contexts:

$$M : \mathbb{R}^D \rightarrow \mathbb{R}^D \text{ with } M(x) = \text{ResonantProjection}(x, R) \quad (11)$$

where R is a set of stored resonance patterns.

The memory module influences the processing dynamics through contextual modulation:

$$W^{\text{eff}} = W^0 + \Delta_{\Phi}(M(x)) \quad (12)$$

where W^0 represents the base weights and Δ_{Φ} is a learned generator that aligns current activation with prior coherent states.

7.2 Theoretical Properties

The resonant memory system exhibits several important theoretical properties:

Proposition 7. *The entropy of the memory retrieval process $H(M(x))$ is minimized when x aligns with existing resonance patterns in R .*

This property ensures that the memory system naturally reinforces existing knowledge structures while still allowing for the formation of new patterns. Recent work on memory-augmented neural networks [14] has shown that similar principles can enhance learning efficiency and knowledge retention.

Theorem 8 (Memory-Induced Resonance). *When the resonant memory modulates processing weights according to $W^{\text{eff}} = W^0 + \Delta_{\Phi}(M(x))$, the system exhibits increased sensitivity to inputs that align with stored resonance patterns.*

This theorem establishes that the memory system creates a form of priming or sensitization to previously encountered patterns, similar to cognitive priming effects observed in human memory [24].

The entropy-modulated resonant memory also introduces a form of experience-dependent plasticity:

Definition 3. *The resonance learning rule updates memory patterns based on alignment and entropy:*

$$R \leftarrow R \cup \{x\} \text{ if } E(x) < \theta \text{ and } \max_{r \in R} \langle x, r \rangle < \delta \quad (13)$$

where $E(x)$ is the entropy of x and θ, δ are thresholds.

This learning rule ensures that the memory system preferentially stores low-entropy, distinctive patterns, aligning with theories of efficient coding in cognitive and neural systems [1].

8 Manipulation Monad Symbolic Processor

The Manipulation Monad provides a structured formalism for tracking and transforming symbolic states within the PRKM architecture. This subsystem bridges traditional symbolic computation with the resonance-based approach of the overall architecture.

8.1 Definition and Construction

Symbolic states σ are represented as triples:

$$\mathbb{X}(\sigma) = (\sigma, T, E) \quad (14)$$

where σ is the symbolic content, T is a set of transformation rules, and E is the entropy measure.

The Monad operates within the Prime Hilbert Space, coupling symbolic entropy $E(\sigma)$ directly to attention entropy H^t :

$$H^{t+1} = H^t + \Delta E(\sigma) \quad (15)$$

Symbolic collapse identified by the Monad explicitly informs the observer-conditioned collapse vector:

$$o_{\text{new}} = o + \gamma \cdot \mathbb{X}(\sigma)_{\text{collapse}} \quad (16)$$

8.2 Theoretical Properties

The Manipulation Monad establishes a unified resonance condition:

$$R(\sigma) \geq R_{\text{threshold}}, E(\sigma) \in [0.2, 0.3], \text{Parity}(\sigma) = \text{even} \quad (17)$$

This condition identifies stable symbolic states that exhibit both resonance and controlled entropy levels. Recent work on symbolic reasoning in neural networks [11] has shown that such hybrid approaches can combine the strengths of symbolic and neural computation.

Theorem 9 (Symbolic-Neural Alignment). *When the Manipulation Monad identifies a symbolic collapse state, the corresponding neural representation in the Prime Hilbert space exhibits localization around prime-structured dimensions.*

This theorem establishes a formal connection between symbolic computation and neural representation in the PRKM architecture, addressing a long-standing challenge in artificial intelligence research [19].

9 Collapse Convergence Theorem

We now present a key theoretical result that establishes the convergence properties of the PRKM architecture. This theorem provides theoretical guarantees about the system's behavior and learning dynamics.

Theorem 10 (Collapse Convergence). *Let R be the full PRKM model. If:*

- *Inputs are separable in H_{prime}*
- *Attention entropy $H^{(t)}$ is decreasing*
- *Observer alignment $\langle x, o \rangle$ is monotonic*

Then the collapse $x^{(t)} \rightarrow x^$ converges in $t < T$ steps, bounded by phase sharpening δ .*

Proof Sketch. Under the given conditions, the entropy reduction dynamics ensure that the attention distribution progressively concentrates on a subset of connections. The monotonic observer alignment ensures that this concentration moves toward states that maximize alignment with the observer vector. The separability condition ensures that different inputs remain distinguishable throughout this process. Together, these conditions establish a contraction mapping on the state space, guaranteeing convergence to a fixed point x^* in finite time. $\square$

An important corollary of this theorem relates to the special properties of tokens with specific harmonic properties:

Corollary 11. *Tokens with digital root 9 and balanced prime exponents converge faster, indicating preferred symbolic collapse into low-entropy states.*

This corollary establishes a theoretical explanation for the empirical observation that certain token patterns exhibit more rapid and stable convergence behavior. Recent work on convergence in attention-based models [13] has identified similar properties, though without the explicit connection to number-theoretic structures.

The collapse convergence theorem provides a theoretical foundation for understanding how the PRKM architecture processes information and forms stable knowledge representations. It also suggests potential optimizations for practical implementations by identifying conditions that facilitate rapid convergence.

10 Resonant Knowledge Learning Objective

The PRKM architecture employs a specialized learning objective that incorporates multiple components aligned with the system’s resonance-based approach to knowledge representation.

10.1 Compound Loss Function

The system minimizes a compound loss function:

$$L = L_{CE} + \lambda_1 S_p + \lambda_2(\text{mod9 Phase Dispersion}) + \lambda_3(1 - \langle x, o \rangle) + \lambda_4 E(\sigma) \quad (18)$$

where:

- L_{CE} is the standard cross-entropy loss for prediction
- S_p is a sparsity penalty on attention weights
- mod9 Phase Dispersion measures the coherence of phase relationships
- $\langle x, o \rangle$ represents observer alignment
- $E(\sigma)$ is the symbolic entropy from the Manipulation Monad

10.2 Theoretical Properties

The resonant knowledge learning objective exhibits several important theoretical properties:

Proposition 12. *Minimizing the compound loss L leads to representations that maximize observer alignment while maintaining low entropy and coherent phase relationships.*

This property ensures that the learning process naturally favors representations that exhibit resonance and alignment with the observer. Recent work on multi-objective optimization in neural networks [25] has shown that such compound objectives can lead to more robust and interpretable models.

Theorem 13 (Resonance Learning). *Under the resonant knowledge learning objective, the system develops an internal representation that preferentially encodes information at harmonic resonance points in the Prime Hilbert space.*

This theorem establishes that the learning process inherently leverages the structural properties of the Prime Hilbert Embedding and Mod 9 Harmonic Phase Encoding to create efficient knowledge representations. Similar principles have been observed in studies of efficient coding in neural systems [21].

The resonant knowledge learning objective also introduces a form of implicit regularization:

Proposition 14. *The phase dispersion term in the learning objective acts as a regularizer that promotes coherent phase relationships and prevents overfitting to non-resonant patterns.*

This regularization effect contributes to the system’s ability to generalize from limited examples, addressing a key challenge in current machine learning approaches [29].

11 Conclusion and Future Work

This paper has presented a comprehensive formalism for the Phase-Shifted Resonance-Based Knowledge Model Architecture, integrating principles from quantum mechanics, number theory, and cognitive science to create a novel approach to knowledge representation and processing.

11.1 Key Contributions

The major contributions of this work include:

- A formal definition of the Prime Hilbert Embedding space that leverages prime number structures for efficient knowledge representation
- The Mod 9 Harmonic Phase Encoding that introduces cyclical harmonic relationships among tokens
- The Iterative Resonance Attention mechanism with provable entropy reduction properties
- The Observer-Conditioned Collapse formalism that establishes a connection to quantum measurement theory
- The Entropy-Modulated Resonant Memory system that facilitates efficient storage and retrieval of resonant patterns
- The Manipulation Monad that bridges symbolic and neural computation
- The Collapse Convergence Theorem that provides theoretical guarantees about system behavior

Together, these elements form a coherent mathematical framework for knowledge processing that offers several advantages over conventional approaches, including faster learning from limited examples, more interpretable internal representations, and better alignment with cognitive principles.

11.2 Future Work

Several promising directions for future research emerge from this work:

- **Empirical Validation:** Implementing the PRKM architecture in practical applications and comparing its performance to conventional approaches.
- **Theoretical Extensions:** Expanding the formalism to incorporate additional number-theoretic structures and quantum-inspired processing mechanisms.
- **Cognitive Alignment:** Further exploring the connections between the PRKM architecture and theories of human cognition and consciousness.
- **Hardware Implementation:** Investigating specialized hardware designs that can efficiently implement the resonance-based processing approach.
- **Symbolic Oracle Feedback:** Developing mechanisms for integrating external symbolic knowledge sources with the resonance-based learning process.

The Phase-Shifted Resonance-Based Knowledge Model Architecture represents a significant step toward reconciling symbolic and neural approaches to artificial intelligence, offering a mathematically rigorous foundation for systems that learn efficiently while maintaining interpretable internal representations aligned with cognitive principles.

Acknowledgments

The author acknowledges Don Gunter for the foundational inspiration behind the triadic domain framework (Counting, Measuring, Manipulation)

References

- [1] Barlow, H.B. (1961). Possible principles underlying the transformations of sensory messages. In W. A. Rosenblith (Ed.), *Sensory Communication* (pp. 217-234). MIT Press.
- [2] Ba, J.L., Kiros, J.R., & Hinton, G.E. (2016). Layer normalization. *arXiv preprint arXiv:1607.06450*.
- [3] Baars, B.J. (2005). Global workspace theory of consciousness: toward a cognitive neuroscience of human experience. *Progress in brain research*, 150, 45-53.
- [4] Bengio, Y., Courville, A., & Vincent, P. (2013). Representation learning: A review and new perspectives. *IEEE transactions on pattern analysis and machine intelligence*, 35(8), 1798-1828.
- [5] Bost, J.B., & Connes, A. (1995). Hecke algebras, type III factors and phase transitions with spontaneous symmetry breaking in number theory. *Selecta Mathematica*, 1(3), 411-457.
- [6] Cohen, T.S., Weiler, M., Kicanaoglu, B., & Welling, M. (2020). Gauge equivariant convolutional networks and the icosahedral CNN. In *International Conference on Machine Learning* (pp. 1321-1330).
- [7] Conway, J.H., & Guy, R.K. (1996). *The book of numbers*. Springer Science & Business Media.

- [8] Dehaene, S. (2011). The number sense: How the mind creates mathematics. Oxford University Press.
- [9] Elhage, N., Nanda, N., Olsson, C., Henighan, T., Joseph, N., Mann, B., ... & Amodei, D. (2021). A mathematical framework for transformer circuits. Transformer Circuits Thread.
- [10] Friston, K. (2010). The free-energy principle: a unified brain theory? *Nature reviews neuroscience*, 11(2), 127-138.
- [11] Garcez, A.D., Gori, M., Lamb, L.C., Serafini, L., Spranger, M., & Tran, S.N. (2019). Neural-symbolic computing: An effective methodology for principled integration of machine learning and reasoning. arXiv preprint arXiv:1905.06088.
- [12] Gobet, F., Lane, P.C., Croker, S., Cheng, P.C., Jones, G., Oliver, I., & Pine, J.M. (2001). Chunking mechanisms in human learning. *Trends in cognitive sciences*, 5(6), 236-243.
- [13] Gong, C., He, D., Tan, X., Qin, T., Wang, L., & Liu, T.Y. (2019). FRAGE: Frequency-agnostic word representation. In *Advances in neural information processing systems* (pp. 1334-1345).
- [14] Graves, A., Wayne, G., Reynolds, M., Harley, T., Danihelka, I., Grabska-Barwińska, A., ... & Hassabis, D. (2016). Hybrid computing using a neural network with dynamic external memory. *Nature*, 538(7626), 471-476.
- [15] Hardy, G.H., & Wright, E.M. (2008). An introduction to the theory of numbers. Oxford university press.
- [16] Lakoff, G., & Núñez, R.E. (1999). Where mathematics comes from. Basic Books.
- [17] Leike, J., Krueger, D., Everitt, T., Martic, M., Maini, V., & Legg, S. (2018). Scalable agent alignment via reward modeling: a research direction. arXiv preprint arXiv:1811.07871.
- [18] Mallat, S. (1999). A wavelet tour of signal processing. Elsevier.
- [19] Marcus, G., & Davis, E. (2020). Rebooting AI: Building artificial intelligence we can trust. Vintage.
- [20] Nielsen, M.A., & Chuang, I. (2010). Quantum computation and quantum information. Cambridge University Press.
- [21] Olshausen, B.A., & Field, D.J. (1996). Emergence of simple-cell receptive field properties by learning a sparse code for natural images. *Nature*, 381(6583), 607-609.
- [22] Preskill, J. (1998). Quantum information and computation. Lecture notes for physics 229.
- [23] Russell, S.J., & Norvig, P. (2010). Artificial intelligence: a modern approach. Pearson Education Limited.
- [24] Schacter, D.L., & Addis, D.R. (2007). The cognitive neuroscience of constructive memory: remembering the past and imagining the future. *Philosophical Transactions of the Royal Society B: Biological Sciences*, 362(1481), 773-786.
- [25] Sener, O., & Koltun, V. (2018). Multi-task learning as multi-objective optimization. In *Advances in Neural Information Processing Systems* (pp. 527-538).
- [26] Sussillo, D., & Barak, O. (2013). Opening the black box: low-dimensional dynamics in high-dimensional recurrent neural networks. *Neural computation*, 25(3), 626-649.
- [27] Tononi, G., Boly, M., Massimini, M., & Koch, C. (2016). Integrated information theory: from consciousness to its physical substrate. *Nature Reviews Neuroscience*, 17(7), 450-461.
- [28] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., ... & Polosukhin, I. (2017). Attention is all you need. In *Advances in neural information processing systems* (pp. 5998-6008).
- [29] Zhang, C., Bengio, S., Hardt, M., Recht, B., & Vinyals, O. (2021). Understanding deep learning (still) requires rethinking generalization. *Communications of the ACM*, 64(3), 107-115.
- [30] Zurek, W.H. (2003). Decoherence, einselection, and the quantum origins of the classical. *Reviews of modern physics*, 75(3), 715.
- [31] Anthropic (2021). A mathematical framework for understanding attention in transformers. Technical blog post.